

Machine Learning Lecture 10

Reinforcement Learning

Qian Ma

Sun Yat-sen University

Spring Semester, 2020



Acknowledgement

- A large part of slides in this lecture are originally from
 - Prof. Andrew Ng (Stanford University)
 - Prof. Weinan Zhang (Shanghai Jiao Tong University)



Prof. Andrew Ng
Stanford University



Prof. Weinan Zhang
Shanghai Jiao Tong University

Content


- Introduction to Reinforcement Learning
- Model-based Reinforcement Learning
 - Markov Decision Process
 - Planning by Dynamic Programming
- Model-free Reinforcement Learning
 - Monte-Carlo
 - Temporal Difference

Content

- Introduction to Reinforcement Learning
- Model-based Reinforcement Learning
 - Markov Decision Process
 - Planning by Dynamic Programming
- Model-free Reinforcement Learning
 - Monte-Carlo
 - Temporal Difference


Reinforcement Learning

- Learning from interaction
 - Given the current situation, what to do next in order to maximize utility?



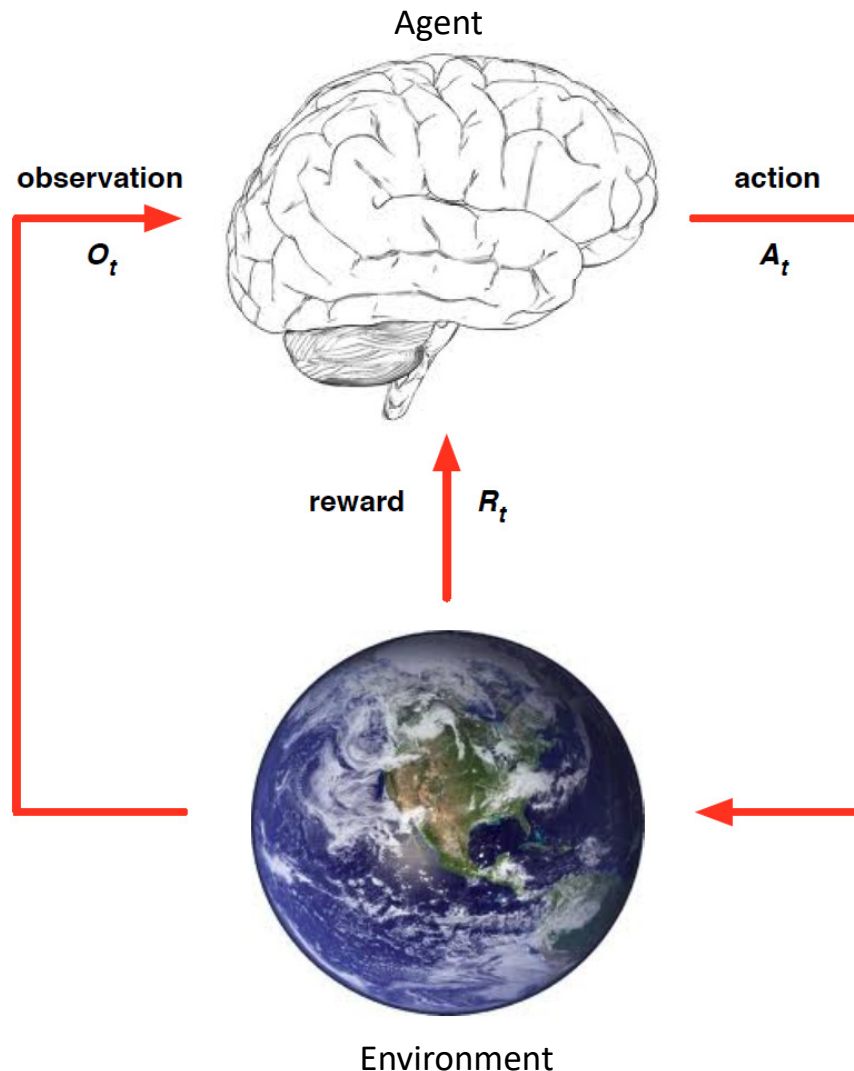
Reinforcement Learning Definition

- A computational approach by learning from interaction to achieve a goal



- Three aspects
 - Sensation: sense the state of the environment to some extent
 - Action: able to take actions that affect the state and achieve the goal
 - Goal: maximize the cumulative reward over time

Reinforcement Learning



- At each step t , the agent
 - Receives observation O_t
 - Receives scalar reward R_t
 - Executes action A_t
- The environment
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits scalar reward R_{t+1}
- t increments at environment step

Elements of RL Systems

- **History** is the sequence of observations, action, rewards

$$H_t = O_1, R_1, A_1, O_2, R_2, A_2, \dots, O_{t-1}, R_{t-1}, A_{t-1}, O_t, R_t$$

- i.e. all observable variables up to time t
- E.g., the sensorimotor stream of a robot or embodied agent
- What happens next depends on the history:
 - The agent selects actions
 - The environment selects observations/rewards
- **State** is the information used to determine what happens next (actions, observations, rewards)
- Formally, state is a function of the history

$$S_t = f(H_t)$$

Elements of RL Systems

- **Policy** is the learning agent's way of behaving at a given time
 - It is a map from state to action
 - Deterministic policy

$$a = \pi(s)$$

- Stochastic policy

$$\pi(a|s) = P(A_t = a|S_t = s)$$

Elements of RL Systems

- Reward
 - A scalar defining the goal in an RL problem
 - For immediate sense of what is good
- Value function
 - State value is a scalar specifying what is good in the long run
 - Value function is a prediction of the cumulative future reward
 - Used to evaluate the goodness/badness of states (given the current policy)

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$

Elements of RL Systems

- Reward

- A scalar defining the goal in an RL problem
- For immediate sense of what is good

- Value function

- State value is a scalar specifying what is a good state in the long run, i.e., the cumulative reward

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$

- Action value is a scalar specifying what is a good action at a specific state in the long run

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a]$$


Elements of RL Systems

- A **Model** of the environment that mimics the behavior of the environment
 - Predict the next state


$$\mathcal{P}_{sa}(s') = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

- Predicts the next (immediate) reward

$$\mathcal{R}_s(a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$




Maze Example




- State: agent's location
- Action: N,E,S,W

Maze Example




- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
 - No move if the action is to the wall

Maze Example



- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step


Maze Example



- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step

- Given a policy as shown above
 - Arrows represent policy $\pi(s)$ for each state s

Maze Example



- State: agent's location
- Action: N,E,S,W
- State transition: move to the next grid according to the action
- Reward: -1 per time step

- Numbers represent value $v_{\pi}(s)$ of each state s

Categorizing RL Agents

- Model based RL
 - Policy and/or value function
 - Model of the environment
 - E.g., the maze game above, game of Go
- Model-free RL
 - Policy and/or value function
 - No model of the environment
 - E.g., general playing Atari games

Content

- Introduction to Reinforcement Learning
- Model-based Reinforcement Learning
 - Markov Decision Process
 - Planning by Dynamic Programming
- Model-free Reinforcement Learning
 - Monte-Carlo
 - Temporal Difference

Markov Property

“The future is independent of the past given the present”

- Definition

- A state S_t is Markov if and only if

$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, \dots, S_t]$$

- Properties

- The state captures all relevant information from the history
 - Once the state is known, the history may be thrown away
 - i.e. the state is sufficient statistic of the future

Markov Decision Process

- A Markov decision process is a tuple $(S, A, \{P_{sa}\}, \gamma, R)$
- S is the set of states
 - E.g., location in a maze, or current screen in an Atari game
- A is the set of actions
 - E.g., move N, E, S, W, or the direction of the joystick and the buttons
- P_{sa} are the state transition probabilities
 - For each state $s \in S$ and action $a \in A$, P_{sa} is a distribution over the next state in S
- $\gamma \in [0,1]$ is the discount factor for the future reward
- $R : S \times A \mapsto \mathbb{R}$ is the reward function
 - Sometimes the reward is only assigned to state

Markov Decision Process

The dynamics of an MDP proceeds as

- Start in a state s_0
- The agent chooses some action $a_0 \in A$
- The agent gets the reward $R(s_0, a_0)$
- MDP randomly transits to some successor state $s_1 \sim P_{s_0 a_0}$
- This proceeds iteratively

$$s_0 \xrightarrow[R(s_0, a_0)]{a_0} s_1 \xrightarrow[R(s_1, a_1)]{a_1} s_2 \xrightarrow[R(s_2, a_2)]{a_2} s_3 \cdots$$

- Until a terminal state s_T or proceeds with no end
- The total payoff of the agent is

$$R(s_0, a_0) + \gamma R(s_1, a_1) + \gamma^2 R(s_2, a_2) + \cdots$$

Reward on State Only

- For a large part of cases, reward is only assigned to the state
 - E.g., in maze game, the reward is on the location
 - In game of Go, the reward is only based on the final territory
- The reward function $R(s) : S \mapsto \mathbb{R}$
- MDPs proceed

$$s_0 \xrightarrow[R(s_0)]{a_0} s_1 \xrightarrow[R(s_1)]{a_1} s_2 \xrightarrow[R(s_2)]{a_2} s_3 \cdots$$

- cumulative reward (total payoff)

$$R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots$$

MDP Goal and Policy

- The goal is to choose actions over time to maximize the expected cumulative reward

$$\mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots]$$

- $\gamma \in [0,1]$ is the discount factor for the future reward, which makes the agent prefer immediate reward to future reward
 - In finance case, today's \$1 is more valuable than \$1 in tomorrow
- Given a particular policy $\pi(s) : S \mapsto A$
 - i.e. take the action $a = \pi(s)$ at state s
- Define the value function for π

$$V^\pi(s) = \mathbb{E}[R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \cdots | s_0 = s, \pi]$$

- i.e. expected cumulative reward given the start state and taking actions according to π

Bellman Equation for Value Function

- Define the value function for π

$$V^\pi(s) = \mathbb{E}[R(s_0) + \underbrace{\gamma R(s_1) + \gamma^2 R(s_2) + \cdots}_{\gamma V^\pi(s_1)} | s_0 = s, \pi]$$

$$= R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s')$$

Bellman Equation

Immediate
Reward

Time
decay

State
transition

Value of
the next
state

Optimal Value Function

- The optimal value function for each state s is best possible sum of discounted rewards that can be attained by any policy

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

- The Bellman's equation for optimal value function

$$V^*(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V^*(s')$$

- The optimal policy

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^*(s')$$

- For every state s and every policy π

$$V^*(s) = V^{\pi^*}(s) \geq V^{\pi}(s)$$

Value Iteration & Policy Iteration

- Note that the value function and policy are correlated

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P_{s\pi(s)}(s') V^\pi(s')$$

$$\pi(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V^\pi(s')$$

- It is feasible to perform iterative update towards the optimal value function and optimal policy
 - Value iteration
 - Policy iteration

Value Iteration

- For an MDP with finite state and action spaces

$$|S| < \infty, |A| < \infty$$

- Value iteration is performed as

1. For each state s , initialize $V(s) = 0$.

2. Repeat until convergence {

For each state, update

$$V(s) = R(s) + \max_{a \in A} \gamma \sum_{s' \in S} P_{sa}(s') V(s')$$

}

- Note that there is no explicit policy in above calculation

Policy Iteration

- For an MDP with finite state and action spaces

$$|S| < \infty, |A| < \infty$$

- Policy iteration is performed as


1. Initialize π randomly
2. Repeat until convergence {
 - a) Let $V := V^\pi$
 - b) For each state, update

$$\pi(s) = \arg \max_{a \in A} \sum_{s' \in S} P_{sa}(s') V(s')$$


}

- The step of value function update could be time-consuming


Policy Iteration



- Policy evaluation
 - Estimate V^π
 - Iterative policy evaluation
- Policy improvement
 - Generate $\pi' \geq \pi$
 - Greedy policy improvement



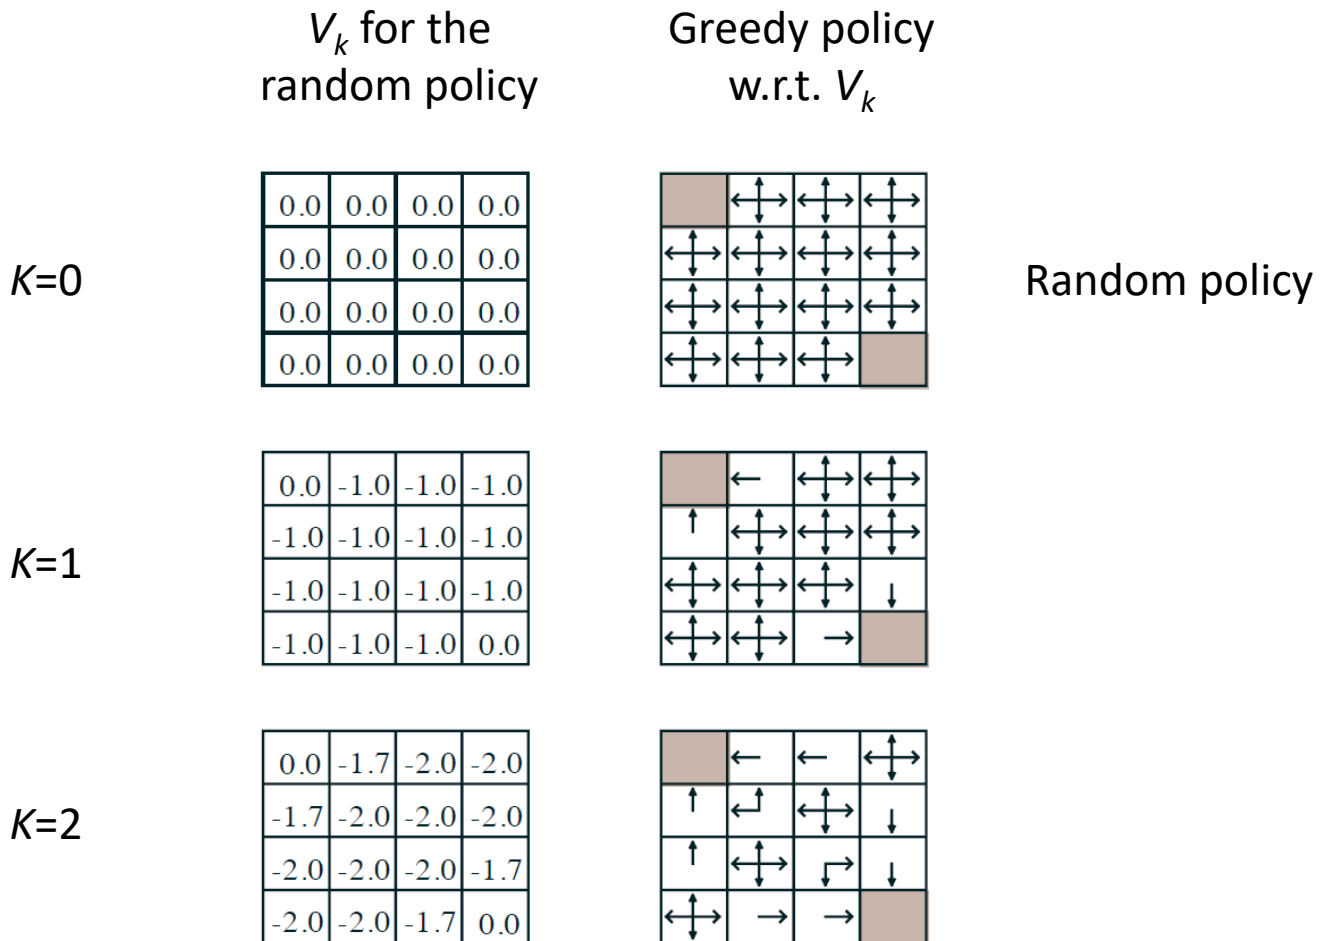
Evaluating a Random Policy in a Small Gridworld



- Undiscounted episodic MDP ($\gamma=1$)
- Nonterminal states 1,...,14
- Two terminal states (shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is -1 until the terminal state is reached
- Agent follows a uniform random policy

$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$$

Evaluating a Random Policy in a Small Gridworld



Evaluating a Random Policy in a Small Gridworld

